

Course: MMDS

Lab 01

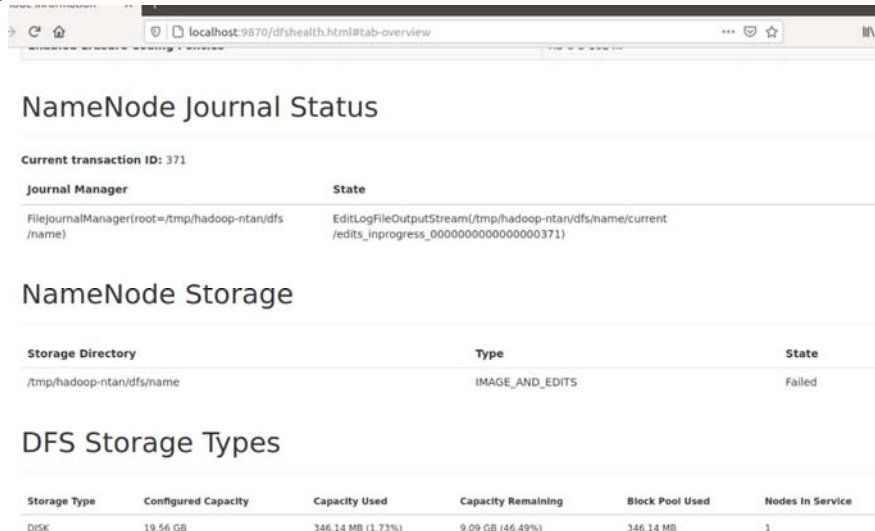
Set up Hadoop

Fill answers of the questions below in the given tables.
Your screenshots must contain commands for required operations.

Question 1:

Set up Hadoop following this [tutorial](#) and then take a screenshot of the hdfs-site ([localhost:9870](#)) which shows information for NameNode Journal Status, NameNode Storage, DFS Storage Types (Overview page).

For example,



The screenshot shows the Hadoop DFS Health page in a web browser. The address bar displays 'localhost:9870/dfshealth.html#tab-overview'. The page is divided into three main sections: NameNode Journal Status, NameNode Storage, and DFS Storage Types.

NameNode Journal Status

Current transaction ID: 371

Journal Manager	State
FileJournalManager(root=/tmp/hadoop-ntan/dfs/name)	EditLogOutputStream(/tmp/hadoop-ntan/dfs/name/current/edits_inprogress_0000000000000000371)

NameNode Storage

Storage Directory	Type	State
/tmp/hadoop-ntan/dfs/name	IMAGE_AND_EDITS	Failed

DFS Storage Types

Storage Type	Configured Capacity	Capacity Used	Capacity Remaining	Block Pool Used	Nodes In Service
DISK	19.56 GB	346.14 MB (1.73%)	9.09 GB (46.49%)	346.14 MB	1

```
libfreerdp-client2-2 libfreerdp2-2 libgc1c2 libgee-0.8-2 libgexiv2-2
libgom-1.0-0 libgpgmepp6 libgpod-common libgpod4 liblangtag-common
liblangtag1 liblirc-client0 liblua5.3-0 libmediaart-2.0-0 libmsspub-0.1-1
libodfgen-0.1-1 libqqwing2v5 libraw16 librevenge-0.0-0 libsgutils2-2
libssh-4 libsuitesparseconfig5 libvncclient1 libwinpr2-2 libxapian30
libxmlsec1-nss lp-solve media-player-info python3-mako python3-markupsafe
syslinux syslinux-common syslinux-legacy usb-creator-common
Use 'sudo apt autoremove' to remove them.
The following additional packages will be installed:
  ca-certificates-java java-common
Suggested packages:
  default-jre fonts-dejavu-extra fonts-ipafont-gothic fonts-ipafont-mincho
  fonts-wqy-microhei fonts-wqy-zenhei
The following NEW packages will be installed:
  ca-certificates-java java-common openjdk-8-jre-headless
0 upgraded, 3 newly installed, 0 to remove and 0 not upgraded.
Need to get 28.3 MB of archives.
After this operation, 104 MB of additional disk space will be used.
Do you want to continue? [Y/n] Y
Get:1 http://us.archive.ubuntu.com/ubuntu bionic-updates/main amd64 java-common
all 0.68ubuntu1~18.04.1 [14.5 kB]
Get:2 http://us.archive.ubuntu.com/ubuntu bionic-updates/universe amd64 openjdk-
8-jre-headless amd64 8u342-b07-0ubuntu1~18.04 [28.3 MB]
63% [2 openjdk-8-jre-headless 20.0 MB/28.3 MB 71%] 262 kB/s 31s
```

```
thanhlam@ubuntu: ~
File Edit View Search Terminal Help
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/rmid to pro
vide /usr/bin/rmid (rmid) in auto mode
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/java to pro
vide /usr/bin/java (java) in auto mode
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/keytool to
provide /usr/bin/keytool (keytool) in auto mode
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/jjs to prov
ide /usr/bin/jjs (jjs) in auto mode
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/pack200 to
provide /usr/bin/pack200 (pack200) in auto mode
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/rmiregistry
to provide /usr/bin/rmiregistry (rmiregistry) in auto mode
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/unpack200 t
o provide /usr/bin/unpack200 (unpack200) in auto mode
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/orbd to pro
vide /usr/bin/orbd (orbd) in auto mode
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/servertool
to provide /usr/bin/servertool (servertool) in auto mode
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/tnameserv t
o provide /usr/bin/tnameserv (tnameserv) in auto mode
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/lib/jexec to pr
ovide /usr/bin/jexec (jexec) in auto mode
Progress: [ 94%] [#####.....]
```



```
thanhlam@ubuntu: ~  
File Edit View Search Terminal Help  
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/java to provide /usr/bin/java (java) in auto mode  
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/keytool to provide /usr/bin/keytool (keytool) in auto mode  
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/jjs to provide /usr/bin/jjs (jjs) in auto mode  
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/pack200 to provide /usr/bin/pack200 (pack200) in auto mode  
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/rmiregistry to provide /usr/bin/rmiregistry (rmiregistry) in auto mode  
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/unpack200 to provide /usr/bin/unpack200 (unpack200) in auto mode  
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/orbd to provide /usr/bin/orbd (orbd) in auto mode  
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/servertool to provide /usr/bin/servertool (servertool) in auto mode  
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/tnameserv to provide /usr/bin/tnameserv (tnameserv) in auto mode  
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/lib/jexec to provide /usr/bin/jexec (jexec) in auto mode  
thanhlam@ubuntu:~$
```

```
thanhlam@ubuntu: ~  
File Edit View Search Terminal Help  
/usr/bin/jps (jps) in auto mode  
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/bin/extcheck to provide /usr/bin/extcheck (extcheck) in auto mode  
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/bin/jarsigner to provide /usr/bin/jarsigner (jarsigner) in auto mode  
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/bin/jmap to provide /usr/bin/jmap (jmap) in auto mode  
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/bin/jstatd to provide /usr/bin/jstatd (jstatd) in auto mode  
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/bin/jhat to provide /usr/bin/jhat (jhat) in auto mode  
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/bin/jdb to provide /usr/bin/jdb (jdb) in auto mode  
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/bin/serialver to provide /usr/bin/serialver (serialver) in auto mode  
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/bin/jfr to provide /usr/bin/jfr (jfr) in auto mode  
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/bin/wsgen to provide /usr/bin/wsgen (wsgen) in auto mode  
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/bin/jcmd to provide /usr/bin/jcmd (jcmd) in auto mode  
thanhlam@ubuntu:~$
```

```
thanhlam@ubuntu: ~  
File Edit View Search Terminal Help  
Setting up ncurses-term (6.1-1ubuntu1.18.04) ...  
Setting up openssh-sftp-server (1:7.6p1-4ubuntu0.7) ...  
Setting up ssh-import-id (5.7-0ubuntu1.1) ...  
Setting up openssh-server (1:7.6p1-4ubuntu0.7) ...  
  
Creating config file /etc/ssh/sshd_config with new version  
Creating SSH2 RSA key; this may take some time ...  
2048 SHA256:SmnmJ0CIyINymHCRA6tyYbw+Fa/EZNT1qVGihPH6X+o root@ubuntu (RSA)  
Creating SSH2 ECDSA key; this may take some time ...  
256 SHA256:xLi20FCQivAPySFusCz4bFUI7g1ybE0d6YD1GPxGfr8 root@ubuntu (ECDSA)  
Creating SSH2 ED25519 key; this may take some time ...  
256 SHA256:B6BZn1ytLRqfu1DHddJvVa+0saNI/BBHxNPGD84EQHw root@ubuntu (ED25519)  
Created symlink /etc/systemd/system/sshd.service → /lib/systemd/system/ssh.service.  
Created symlink /etc/systemd/system/multi-user.target.wants/ssh.service → /lib/systemd/system/ssh.service.  
Setting up ssh (1:7.6p1-4ubuntu0.7) ...  
Processing triggers for man-db (2.8.3-2ubuntu0.1) ...  
Processing triggers for ufw (0.36-0ubuntu0.18.04.2) ...  
Processing triggers for ureadahead (0.100.0-21) ...  
Processing triggers for systemd (237-3ubuntu10.53) ...  
thanhlam@ubuntu:~$  
  
thanhlam@ubuntu: ~  
File Edit View Search Terminal Help  
Get:2 http://us.archive.ubuntu.com/ubuntu bionic/universe amd64 genders amd64 1.21-1build5 [30.9 kB]  
Get:3 http://us.archive.ubuntu.com/ubuntu bionic/universe amd64 pdsh amd64 2.31-3build2 [108 kB]  
Fetched 170 kB in 3s (61.6 kB/s)  
Preconfiguring packages ...  
Selecting previously unselected package libgenders0:amd64.  
(Reading database ... 154406 files and directories currently installed.)  
Preparing to unpack .../libgenders0_1.21-1build5_amd64.deb ...  
Unpacking libgenders0:amd64 (1.21-1build5) ...  
Selecting previously unselected package genders.  
Preparing to unpack .../genders_1.21-1build5_amd64.deb ...  
Unpacking genders (1.21-1build5) ...  
Selecting previously unselected package pdsh.  
Preparing to unpack .../pdsh_2.31-3build2_amd64.deb ...  
Unpacking pdsh (2.31-3build2) ...  
Setting up libgenders0:amd64 (1.21-1build5) ...  
Setting up genders (1.21-1build5) ...  
Setting up pdsh (2.31-3build2) ...  
Processing triggers for man-db (2.8.3-2ubuntu0.1) ...  
Processing triggers for libc-bin (2.27-3ubuntu1.6) ...  
thanhlam@ubuntu:~$
```



```

thanhlam@ubuntu:~$ ssh-keygen -t rsa -P '' -f ~/.ssh/id_rsa
Generating public/private rsa key pair.
Created directory '/home/thanhlam/.ssh'.
Your identification has been saved in /home/thanhlam/.ssh/id_rsa.
Your public key has been saved in /home/thanhlam/.ssh/id_rsa.pub.
The key fingerprint is:
SHA256:btkdKosoKIQ1uRpvGsam70oeCao4oft5CTmR6Uj2XC0 thanhlam@ubuntu
The key's randomart image is:
+----[RSA 2048]-----+
|
|   +   .
|  oB   E .
| =+o= . . S .
| O.B0   . o o .
| =& o . = o .
| #.=.o. o o
| O@=o. . .
+----[SHA256]-----+
thanhlam@ubuntu:~$ cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
thanhlam@ubuntu:~$ chmod 0600 ~/.ssh/authorized_keys
thanhlam@ubuntu:~$

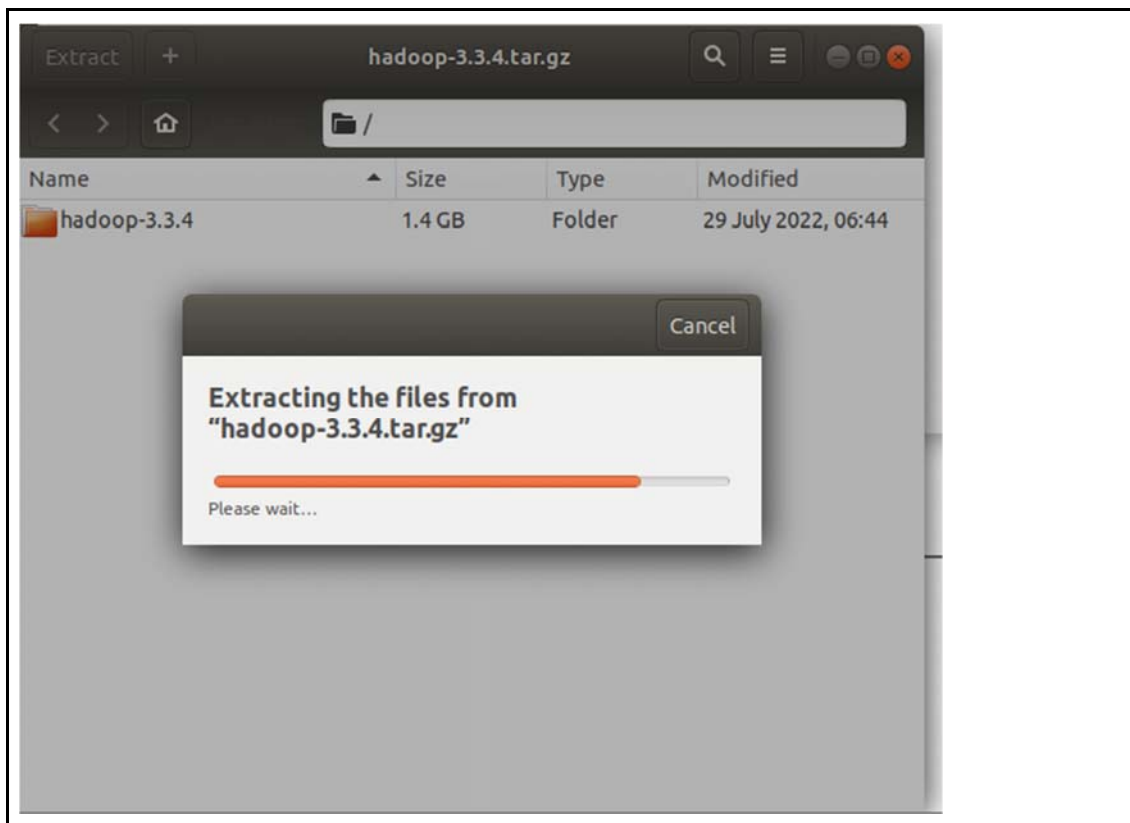
```

```

thanhlam@ubuntu: ~
File Edit View Search Terminal Help
libgom-1.0-0 libpgmepp6 libgpod-common libgpod4 liblangtag-common
liblangtag1 liblirc-client0 liblua5.3-0 libmediaart-2.0-0 libmspub-0.1-1
libodfgen-0.1-1 libqqwing2v5 libraw16 librevenge-0.0-0 libsgutils2-2
libssh-4 libsuitesparseconfig5 libvncclient1 libwinpr2-2 libxapian30
libxmlsec1-nss lp-solve media-player-info python3-mako python3-markupsafe
syslinux syslinux-common syslinux-legacy usb-creator-common
Use 'sudo apt autoremove' to remove them.
The following NEW packages will be installed:
gpp
0 upgraded, 1 newly installed, 0 to remove and 0 not upgraded.
Need to get 46.7 kB of archives.
After this operation, 99.3 kB of additional disk space will be used.
Get:1 http://us.archive.ubuntu.com/ubuntu bionic/universe amd64 gpp amd64 2.24-3
build1 [46.7 kB]
Fetched 46.7 kB in 2s (24.1 kB/s)
Selecting previously unselected package gpp.
(Reading database ... 154448 files and directories currently installed.)
Preparing to unpack .../gpp_2.24-3build1_amd64.deb ...
Unpacking gpp (2.24-3build1) ...
Setting up gpp (2.24-3build1) ...
Processing triggers for man-db (2.8.3-2ubuntu0.1) ...
thanhlam@ubuntu:~$

```

```
thanhlam@ubuntu: ~  
File Edit View Search Terminal Help  
-U : user-defined syntax (specified in 9 following args; see manual)  
-M : user-defined syntax for meta-macros (specified in 7 following args)  
  
-o : output to outfile  
-O : output to outfile and stdout  
-z : line terminator is CR-LF (MS-DOS style)  
-x : enable #exec built-in macro  
-m : enable automatic mode switching upon including .h/.c files  
-n : send LF characters serving as macro terminators to output  
+c : use next 2 args as comment start and comment end sequences  
+s : use next 3 args as string start, end and quote character  
  
Long options:  
--include file : process file before infile  
--nostdinc : don't search standard directories for files to include  
--nocurinc : don't search the current directory for files to include  
--curdirinclast : search the current directory last  
--warninglevel n : set warning level  
--includemarker formatstring : keep track of #include directives in output  
  
--version : display version information and exit  
-h, --help : display this message and exit  
  
thanhlam@ubuntu:~$  
libodfgen-0.1-1 libqqwing2v5 libraw16 libvenge-0.0-0 libsgutils2-2  
libssh-4 libsuitesparseconfig5 libvncclient1 libwinpr2-2 libxapian30  
libxmlsec1-nss lp-solve media-player-info python3-mako python3-markupsafe  
syslinux syslinux-common syslinux-legacy usb-creator-common  
Use 'sudo apt autoremove' to remove them.  
The following NEW packages will be installed:  
  pgpgpg  
0 upgraded, 1 newly installed, 0 to remove and 0 not upgraded.  
Need to get 14.9 kB of archives.  
After this operation, 48.1 kB of additional disk space will be used.  
Get:1 http://us.archive.ubuntu.com/ubuntu bionic/universe amd64 pgpgpg amd64 0.13-9.1build1 [14.9 kB]  
Fetched 14.9 kB in 2s (8,614 B/s)  
Selecting previously unselected package pgpgpg.  
(Reading database ... 154459 files and directories currently installed.)  
Preparing to unpack .../pgpgpg_0.13-9.1build1_amd64.deb ...  
Unpacking pgpgpg (0.13-9.1build1) ...  
Setting up pgpgpg (0.13-9.1build1) ...  
update-alternatives: using /usr/bin/pgpgpg to provide /usr/bin/pgp (pgp) in auto mode  
Processing triggers for man-db (2.8.3-2ubuntu0.1) ...  
thanhlam@ubuntu:~$
```



```
hadoop-env.sh
~/Desktop/hadoop-3.3.4/etc/hadoop

# export HDFS_DFSROUTER_OPTS=""

###
# HDFS StorageContainerManager specific parameters
###
# Specify the JVM options to be used when starting the HDFS Storage Container Manager.
# These options will be appended to the options specified as HADOOP_OPTS
# and therefore may override any similar flags set in HADOOP_OPTS
#
# export HDFS_STORAGECONTAINERMANAGER_OPTS=""

###
# Advanced Users Only!
###

#
# When building Hadoop, one can add the class paths to the commands
# via this special env var:
# export HADOOP_ENABLE_BUILD_PATHS="true"

#
# To prevent accidents, shell commands be (superficially) locked
# to only allow certain users to execute certain subcommands.
# It uses the format of (command)_(subcommand)_USER.
#
# For example, to limit who can execute the namenode command,
# export HDFS_NAMENODE_USER=hdfs

###
# Registry DNS specific parameters
###
# For privileged registry DNS, user to run as after dropping privileges
# This will replace the hadoop.id.str Java property in secure mode.
# export HADOOP_REGISTRYDNS_SECURE_USER=yarn

# Supplemental options for privileged registry DNS
# By default, Hadoop uses jsvc which needs to know to launch a
# server jvm.
# export HADOOP_REGISTRYDNS_SECURE_EXTRA_OPTS="-jvm server"
export JAVA_HOME=/usr/lib/jvm/java-1.8.0-openjdk-amd64
export PATH=${JAVA_HOME}/bin:${PATH}
export HADOOP_CLASSPATH=${JAVA_HOME}/lib/tools.jar

sh Tab Width: 8 Ln 433, Col 51 INS
```



```
merics/ policy.xml user- site.xml xml sn
thanhlam@ubuntu: ~/Desktop/hadoop-3.3.4
File Edit View Search Terminal Help
dtutil      operations related to delegation tokens
envvars     display computed Hadoop environment variables
fs          run a generic filesystem user client
gridmix     submit a mix of synthetic job, modeling a profiled from
            production load
jar <jar>   run a jar file. NOTE: please use "yarn jar" to launch YARN
            applications, not this command.
jnipath     prints the java.library.path
kdiag       Diagnose Kerberos Problems
kerbname    show auth_to_local principal conversion
key         manage keys via the KeyProvider
rumenfolder scale a rumen input trace
rumentrace  convert logs into a rumen trace
s3guard     manage metadata on S3
trace       view and modify Hadoop tracing settings
version     print the version

Daemon Commands:

kms         run KMS, the Key Management Server
registrydns run the registry DNS server

SUBCOMMAND may print help when invoked w/o parameters or with -h.
thanhlam@ubuntu:~/Desktop/hadoop-3.3.4$

<?xml version="1.0" encoding="UTF-8"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
<!--
  Licensed under the Apache License, Version 2.0 (the "License");
  you may not use this file except in compliance with the License.
  You may obtain a copy of the License at

    http://www.apache.org/licenses/LICENSE-2.0

  Unless required by applicable law or agreed to in writing, software
  distributed under the License is distributed on an "AS IS" BASIS,
  WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
  See the License for the specific language governing permissions and
  limitations under the License. See accompanying LICENSE file.
-->

<!-- Put site-specific property overrides in this file. -->

<configuration>
  <property>
    <name>fs.defaultFS</name>
    <value>hdfs://localhost:9000</value>
  </property>
</configuration>
```

```

<?xml version="1.0" encoding="UTF-8"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
<!--
Licensed under the Apache License, Version 2.0 (the "License");
you may not use this file except in compliance with the License.
You may obtain a copy of the License at

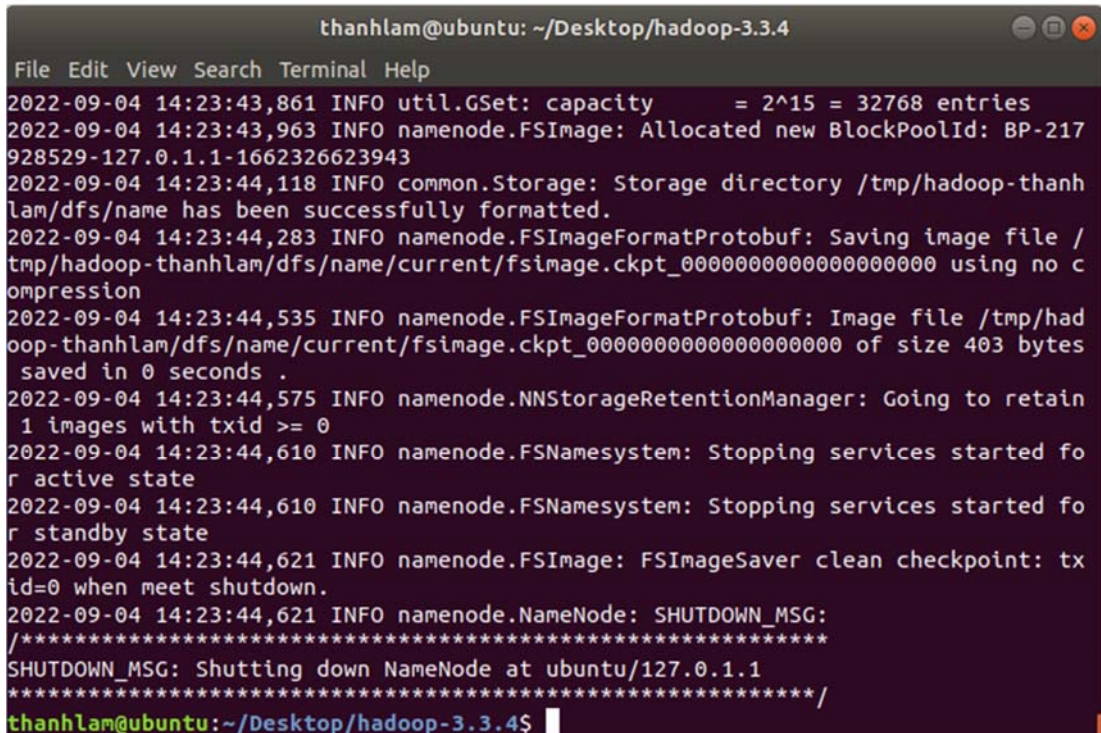
http://www.apache.org/licenses/LICENSE-2.0

Unless required by applicable law or agreed to in writing, software
distributed under the License is distributed on an "AS IS" BASIS,
WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
See the License for the specific language governing permissions and
limitations under the License. See accompanying LICENSE file.
-->

<!-- Put site-specific property overrides in this file. -->

<configuration>
  <property>
    <name>dfs.replication</name>
    <value>1</value>
  </property>
</configuration>

```



The screenshot shows a terminal window titled 'thanhlam@ubuntu: ~/Desktop/hadoop-3.3.4'. The terminal displays a series of log messages from Hadoop's NameNode, including information about GSet capacity, block pool allocation, storage directory formatting, and image saving. The logs conclude with a 'SHUTDOWN_MSG' indicating the NameNode is shutting down at 127.0.1.1. The prompt 'thanhlam@ubuntu:~/Desktop/hadoop-3.3.4\$' is visible at the bottom.

```

thanhlam@ubuntu: ~/Desktop/hadoop-3.3.4
File Edit View Search Terminal Help
2022-09-04 14:23:43,861 INFO util.GSet: capacity      = 2^15 = 32768 entries
2022-09-04 14:23:43,963 INFO namenode.FSImage: Allocated new BlockPoolId: BP-217
928529-127.0.1.1-1662326623943
2022-09-04 14:23:44,118 INFO common.Storage: Storage directory /tmp/hadoop-thanh
lam/dfs/name has been successfully formatted.
2022-09-04 14:23:44,283 INFO namenode.FSImageFormatProtobuf: Saving image file /
tmp/hadoop-thanhlam/dfs/name/current/fsimage.ckpt_000000000000000000 using no c
ompression
2022-09-04 14:23:44,535 INFO namenode.FSImageFormatProtobuf: Image file /tmp/had
oop-thanhlam/dfs/name/current/fsimage.ckpt_000000000000000000 of size 403 bytes
saved in 0 seconds .
2022-09-04 14:23:44,575 INFO namenode.NNStorageRetentionManager: Going to retain
1 images with txid >= 0
2022-09-04 14:23:44,610 INFO namenode.FSNamesystem: Stopping services started fo
r active state
2022-09-04 14:23:44,610 INFO namenode.FSNamesystem: Stopping services started fo
r standby state
2022-09-04 14:23:44,621 INFO namenode.FSImage: FSImageSaver clean checkpoint: tx
id=0 when meet shutdown.
2022-09-04 14:23:44,621 INFO namenode.NameNode: SHUTDOWN_MSG:
/*****
SHUTDOWN_MSG: Shutting down NameNode at ubuntu/127.0.1.1
*****/
thanhlam@ubuntu:~/Desktop/hadoop-3.3.4$

```



```
thanhlam@ubuntu: ~/Desktop/hadoop-3.3.4
File Edit View Search Terminal Help
2022-09-04 14:23:44,283 INFO namenode.FSImageFormatProtobuf: Saving image file /
/tmp/hadoop-thanhlam/dfs/name/current/fsimage.ckpt_00000000000000000000 using no c
ompression
2022-09-04 14:23:44,535 INFO namenode.FSImageFormatProtobuf: Image file /tmp/had
oop-thanhlam/dfs/name/current/fsimage.ckpt_00000000000000000000 of size 403 bytes
saved in 0 seconds .
2022-09-04 14:23:44,575 INFO namenode.NNStorageRetentionManager: Going to retain
1 images with txid >= 0
2022-09-04 14:23:44,610 INFO namenode.FSNamesystem: Stopping services started fo
r active state
2022-09-04 14:23:44,610 INFO namenode.FSNamesystem: Stopping services started fo
r standby state
2022-09-04 14:23:44,621 INFO namenode.FSImage: FSImageSaver clean checkpoint: tx
id=0 when meet shutdown.
2022-09-04 14:23:44,621 INFO namenode.NameNode: SHUTDOWN_MSG:
/*****
SHUTDOWN_MSG: Shutting down NameNode at ubuntu/127.0.1.1
*****/
thanhlam@ubuntu:~/Desktop/hadoop-3.3.4$ sbin/start-dfs.sh
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [ubuntu]
ubuntu: Warning: Permanently added 'ubuntu' (ECDSA) to the list of known hosts.
thanhlam@ubuntu:~/Desktop/hadoop-3.3.4$
```

NameNode Journal Status

Current transaction ID: 3

Journal Manager	State
FileJournalManager(root=/tmp/hadoop-thanhlam/dfs/name)	EditLogFileOutputStream(/tmp/hadoop-thanhlam/dfs/name/current/edits_inprogress_00000000000000000003)

NameNode Storage

Storage Directory	Type	State
/tmp/hadoop-thanhlam/dfs/name	IMAGE_AND_EDITS	Active

DFS Storage Types

Storage Type	Configured Capacity	Capacity Used	Capacity Remaining	Block Pool Used	Nodes In Service
DISK	19.52 GB	24 KB (0%)	9.14 GB (46.82%)	24 KB	1

Hadoop,
2022.


```

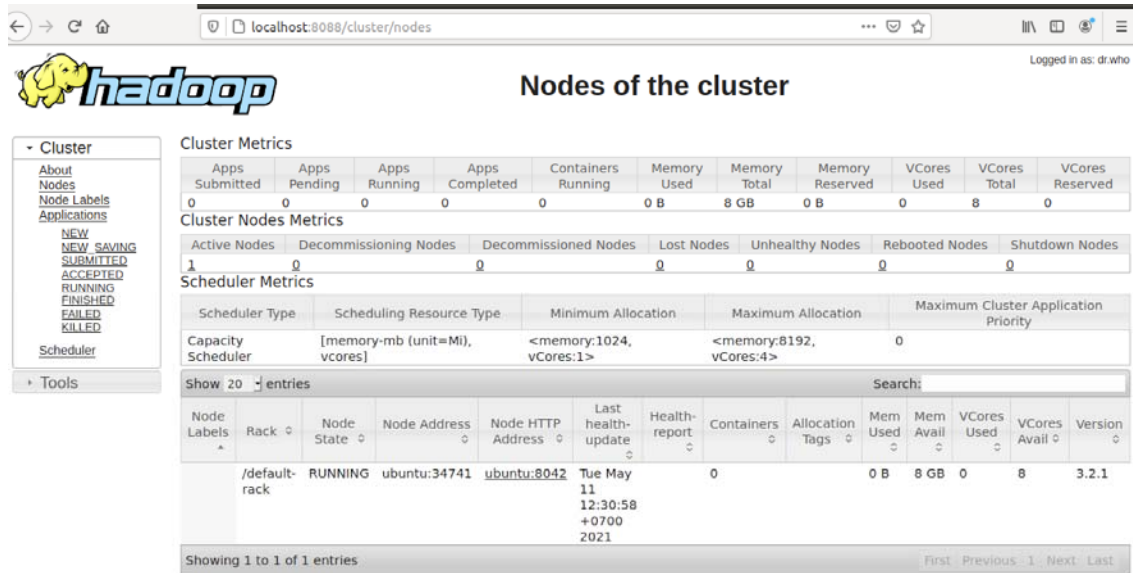
thanhlam@ubuntu:~/Desktop/hadoop-3.3.4$ sbin/stop-dfs.sh
Stopping namenodes on [localhost]
Stopping datanodes
Stopping secondary namenodes [ubuntu]

```

Question 2:

Set up Hadoop following this [tutorial](#) and then take a screenshot of the yarn-site (localhost:8088) which shows information for nodes (<http://localhost:8088/cluster/nodes>).

For example,



The screenshot shows the Hadoop YARN web interface at localhost:8088/cluster/nodes. The page displays various cluster metrics and a list of nodes.

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	VCores Used	VCores Total	VCores Reserved
0	0	0	0	0	0 B	8 GB	0 B	0	8	0

Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes	Rebooted Nodes	Shutdown Nodes
1	0	0	0	0	0	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation	Maximum Cluster Application Priority
Capacity Scheduler	[memory-mb (unit=M), vcores]	<memory:1024, vCores:1>	<memory:8192, vCores:4>	0

Nodes

Node Labels	Rack	Node State	Node Address	Node HTTP Address	Last health-update	Health-report	Containers	Allocation Tags	Mem Used	Mem Avail	VCores Used	VCores Avail	Version
/default-rack		RUNNING	ubuntu:34741	ubuntu:8042	Tue May 11 12:30:58 +0700 2021		0		0 B	8 GB	0	8	3.2.1

Showing 1 to 1 of 1 entries

File Edit View Search Terminal Help

```
    at java.security.AccessController.doPrivileged(Native Method)
    at javax.security.auth.Subject.doAs(Subject.java:422)
    at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1878)
    at org.apache.hadoop.ipc.Server$Handler.run(Server.java:2976)

    at org.apache.hadoop.ipc.Client.getRpcResponse(Client.java:1612)
    at org.apache.hadoop.ipc.Client.call(Client.java:1558)
    at org.apache.hadoop.ipc.Client.call(Client.java:1455)
    at org.apache.hadoop.ipc.ProtobufRpcEngine2$Invoker.invoke(ProtobufRpcEngine2.java:242)
    at org.apache.hadoop.ipc.ProtobufRpcEngine2$Invoker.invoke(ProtobufRpcEngine2.java:129)
    at com.sun.proxy.$Proxy9.addBlock(Unknown Source)
    at org.apache.hadoop.hdfs.protocolPB.ClientNamenodeProtocolTranslatorPB.addBlock(ClientNamenodeProtocolTranslatorPB.java:530)
    at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
    at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
    at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
    at java.lang.reflect.Method.invoke(Method.java:498)
    at org.apache.hadoop.io.retry.RetryInvocationHandler.invokeMethod(RetryInvocationHandler.java:422)
    at org.apache.hadoop.io.retry.RetryInvocationHandler$Call.invokeMethod(RetryInvocationHandler.java:165)
    at org.apache.hadoop.io.retry.RetryInvocationHandler$Call.invoke(RetryInvocationHandler.java:157)
    at org.apache.hadoop.io.retry.RetryInvocationHandler$Call.invokeOnce(RetryInvocationHandler.java:95)
    at org.apache.hadoop.io.retry.RetryInvocationHandler.invoke(RetryInvocationHandler.java:359)
    at com.sun.proxy.$Proxy10.addBlock(Unknown Source)
    at org.apache.hadoop.hdfs.DFSOutputStream.addBlock(DFSOutputStream.java:1088)
    at org.apache.hadoop.hdfs.DataStreamer.locateFollowingBlock(DataStreamer.java:1915)
    at org.apache.hadoop.hdfs.DataStreamer.nextBlockOutputStream(DataStreamer.java:1717)
    at org.apache.hadoop.hdfs.DataStreamer.run(DataStreamer.java:713)
Error: File /user/thanhlam/input/yarn-site.xml._COPYING_ could only be written to 0 of the 1 minReplication nodes. There are 0 datanode(s) running and 0 node(s) are excluded in this operation.
thanhlam@ubuntu:~/Desktop/hadoop-3.3.4$
```

```

<?xml version="1.0"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
<!--
  Licensed under the Apache License, Version 2.0 (the "License");
  you may not use this file except in compliance with the License.
  You may obtain a copy of the License at

    http://www.apache.org/licenses/LICENSE-2.0

  Unless required by applicable law or agreed to in writing, software
  distributed under the License is distributed on an "AS IS" BASIS,
  WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
  See the License for the specific language governing permissions and
  limitations under the License. See accompanying LICENSE file.
-->

<!-- Put site-specific property overrides in this file. -->

<configuration>
  <property>
    <name>mapreduce.framework.name</name>
    <value>yarn</value>
  </property>
  <property>
    <name>mapreduce.application.classpath</name>
    <value>${HADOOP_MAPRED_HOME}/share/hadoop/mapreduce/*:
    $HADOOP_MAPRED_HOME/share/hadoop/mapreduce/lib/*</value>
  </property>
</configuration>

```

```

<?xml version="1.0"?>
<!--
  Licensed under the Apache License, Version 2.0 (the "License");
  you may not use this file except in compliance with the License.
  You may obtain a copy of the License at

    http://www.apache.org/licenses/LICENSE-2.0

  Unless required by applicable law or agreed to in writing, software
  distributed under the License is distributed on an "AS IS" BASIS,
  WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
  See the License for the specific language governing permissions and
  limitations under the License. See accompanying LICENSE file.
-->

<configuration>


<!-- Site specific YARN configuration properties -->
  <property>
    <name>yarn.nodemanager.aux-services</name>
    <value>mapreduce_shuffle</value>
  </property>
  <property>
    <name>yarn.nodemanager.env-whitelist</name>
    <value>JAVA_HOME,HADOOP_COMMON_HOME,HADOOP_HDFS_HOME,HADOOP_CONF_DIR,CLASSPATH_PREPEND_DISTCACHE,HADOOP_YARN_HOME,HADOOP_MAPRED_HOME</value>
  </property>
</configuration>

```

```

thanhlam@ubuntu:~/Desktop/hadoop-3.3.4$ sbin/start-yarn.sh
Starting resourcemanager
Starting nodemanagers

```

All Applications

Cluster

About

Nodes

Node Labels

Applications

NEW

NEW SAVING

SUBMITTED

ACCEPTED

RUNNING

FINISHED

FAILED

KILLED

Scheduler

Tools

Cluster Metrics

Apps Submitted	0	Apps Pending	0	Apps Running	0	Apps Completed	0	Containers Running		Used Resources	<memory:0 B, vCores:0>	<memory:0 B, vCores:0>
----------------	---	--------------	---	--------------	---	----------------	---	--------------------	--	----------------	------------------------	------------------------

Cluster Nodes Metrics

Active Nodes	1	Decommissioning Nodes	0	Decommissioned Nodes	0	Lost Nodes	0
--------------	---	-----------------------	---	----------------------	---	------------	---

Scheduler Metrics

Scheduler Type	Capacity Scheduler	Scheduling Resource Type	[memory-mb (unit=Mi), vCores]	Minimum Allocation	<memory:1024, vCores:1>	Maximum Allocation	<memory:8192, vCores:1>
----------------	--------------------	--------------------------	-------------------------------	--------------------	-------------------------	--------------------	-------------------------

Show 20 entries

ID	User	Name	Application Type	Application Tags	Queue	Application Priority	StartTime	LaunchTime	FinishTime	State	FinalStatus	Running Containers	Allocated Resources
No data available in table													

Showing 0 to 0 of 0 entries

```
thanhlam@ubuntu:~/Desktop/hadoop-3.3.4$ sbin/stop-yarn.sh
Stopping nodemanagers
Stopping resourcemanager
```

Submission Notice

- Export your answer file as pdf
- Rename the pdf following the format:
 <student number>_HoTen.pdf

E.g: 123456_NguyenThanhAn.pdf

If you have not been assigned a student number yet, then use 123456 instead.

- Careless mistakes in filename, format, question order, etc. are not accepted (0 pts).